

Take-home task

# Helping users build trust with Driftmail

Give more control and visibility at the right time

Yuliia Kovalova • Design Lead @SmartBear

# Core Features

Driftmail archives, replies, forwards and unsubscribes on the user's behalf, without approving each action.

It runs in the background. The only thing users see is a weekly digest.

Problem

Users don't trust it enough to adopt it.

- Lack of control - too autonomous
- see what happened (weekly digest?)

- makes not agreed decisions about important things
- needs a way to recover

Driftmail's team has a problem: users don't trust it enough to adopt it.

**In interviews, five users said:**

- "I just want to see what it did. I found out it had auto-archived a bunch of stuff, and I had no way to go back and check what I might have missed. If there was some kind of log I could scroll through, I'd feel a lot better about leaving it on."
- "I'd turn it on if I could watch it first. Like, show me what you're about to do for a few days, and if it looks right, then go ahead and start doing it for real."
- "It's not that I don't want it doing things. Newsletters? Go nuts. But if someone I've never emailed before writes me, I want to see that before anything happens to it."
- "It archived something important and I didn't notice for three days. If I could've just hit undo right away I probably would've kept it on. But there was no way to pull it back."
- "The thing that scares me is it sending a reply as me. I don't care if it sends the idea that it could write something to a client or my boss without me. No. I need to approve anything that goes out with my name on it."

- not showing decisions proactively
- a way to review and confirm

- acts on user behalf without approval
- need for approval

**Some other stakeholder input:**

- "Honestly, the fix is simple, make the onboarding punchier. A slick animation, a progress bar, maybe a confetti moment when they flip on autonomy. Once they see the magic, they'll stick. We're overthinking the trust thing, it's not that important." - Sales manager
- "Half the tickets we get about Driftmail 'doing something wrong' turn out to be totally reasonable calls - the user just didn't know it happened until someone asked them about an email they never saw. The problem isn't accuracy, it's that people find out too late." - Head of support

# Quotes analysis: issues & needs

| Issue                                             | Need                     | Quote                                                                                                                     |
|---------------------------------------------------|--------------------------|---------------------------------------------------------------------------------------------------------------------------|
| Lack of control and visibility                    | See what it did          | “If there was some kind of log I could scroll through, I’d feel a lot better.”                                            |
| Not showing decisions proactively                 | Surface decisions        | “Like, show me what you're about to do for a few days, and if it looks right, then go ahead and start doing it for real.” |
| Makes not agreed decisions about important things | Undo decisions           | “If I could've just hit undo right away I probably would've kept it on. But there was no way to pull it back.”            |
| Acts on user behalf without approval              | Approve high-stake items | "The thing that scares me is it sending a reply as me. ... I need to approve anything that goes out with my name on it.”  |

# Assumptions

## People check email all day

A weekly digest might not be enough to build trust.

**Confirm with users**

## People want the noise gone

And make important stuff to stand out.

**Confirm with users**

## Simple cases are safe to automate

The risk are the mails that similar to them. That's where the user decides.

**Confirm with PM**

## The agent learns from user choices

When people agree with agent, it understands the context better and improves decision making

**Confirm with Engineering**

## Preventing mistakes is better than fixing them

Agent must explain decision and ask for user review when unsure rather than user undoing wrong calls.

**Confirm with PM**

# Core problem

that worth solving

**Lack of visibility and control over  
AI decisions**

Decisions are communicated too late, users  
cannot make a call when needed.

# Solution key pillars

## Show evidence and explain

Show what context agent based replies on

## Help user notice

Show items where agent did it's work

## Approve high-stake decisions

Never sends as you. Only you press Send

## Confirm rather than undo

Each decision teaches agent, trust earned through a visible track record.

Ideas bank, for later

## Activity log

Everything it did today, with Undo

## Trust nudge

It proposes the next step with evidence

## Permissions

control per action: Off, Ask me or Automatic

## Live toast

“Archived 6 newsletters · Undo” as it acts

# Two stories

to cut down the scope

## Story 1 • Needs your review

As a user I want to see the items that need my review in a mailbox folder, so that Driftmail learns from my decisions.

### Addresses:

Catching wrong agent calls early, a need to undo decisions.

## Story 2 • Drafts to approve

As a user I want to see emails where Driftmail drafted a reply, so that I can approve it before it's sent.

### Addresses:

Not having control over the emails user send

# Solution Demo

Interactive prototype

## How will we know it works?

### Driftmail still on at day 30

People continue using, adoption grows



### Manually approved items

Number of items agent required manual review decreased



### “Didn’t know it happened” tickets

Addresses Head of Support’s challenge



### Drafts sent with light edits

Drafts sound like the user, content makes logical sense



## Open Questions

Requires confirming with Product/  
Engineering

**How do we triage importance?**

Only real people's mail become drafts; campaigns get archived? Etc.

**Where does the agent live?**

A Gmail extension, or a scheduled Claude skill? It decides how we handle the AI evidence and output.

**What about mobile?**

How review and approval work when the user is on their phone.

# Thank you

Problem framing and solution prototyping in this project were AI assisted, which sped up the work significantly.